

# Practice Supervised Modelling

Red & White Consulting Partners LLP



# Table of Content

Supervised Learning: Logistic  
Regression



Supervised Learning: Random  
Forest



Technical Details



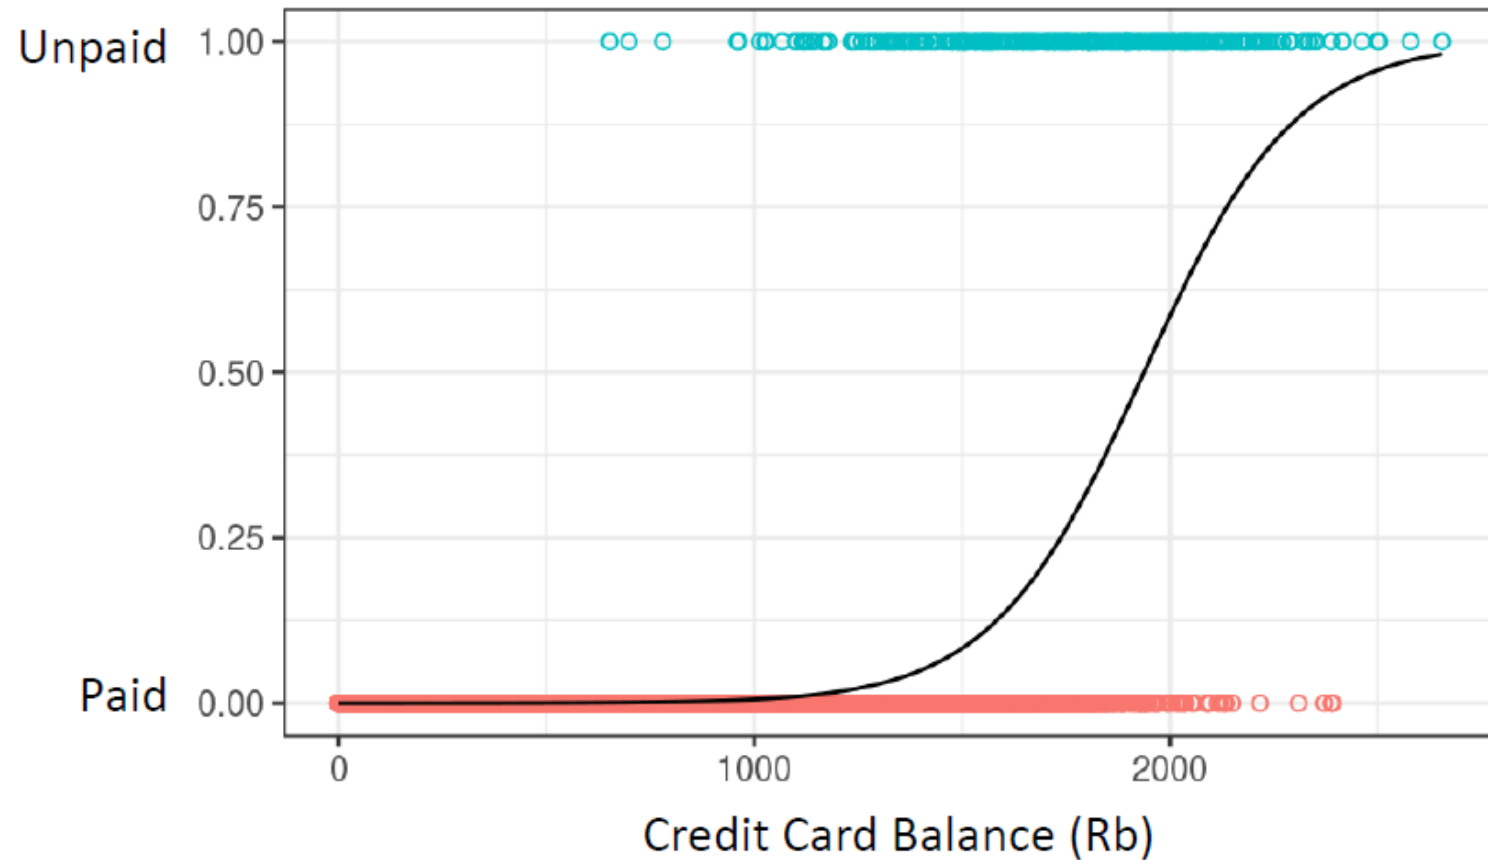
Exercise



# Supervised Learning: Logistic Regression

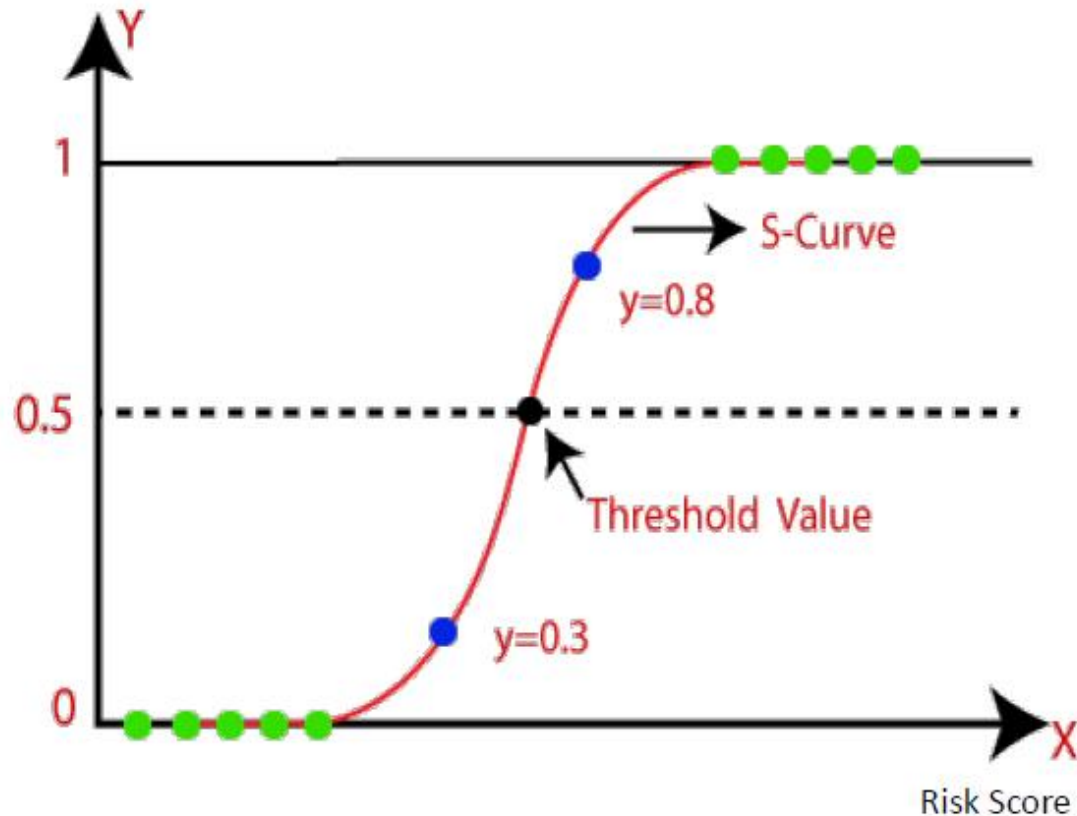
# What is Logistic Regression

Logistic Regression models the probability of the response variable belonging to a specific class



# What is Sigmoid Function

Default Status

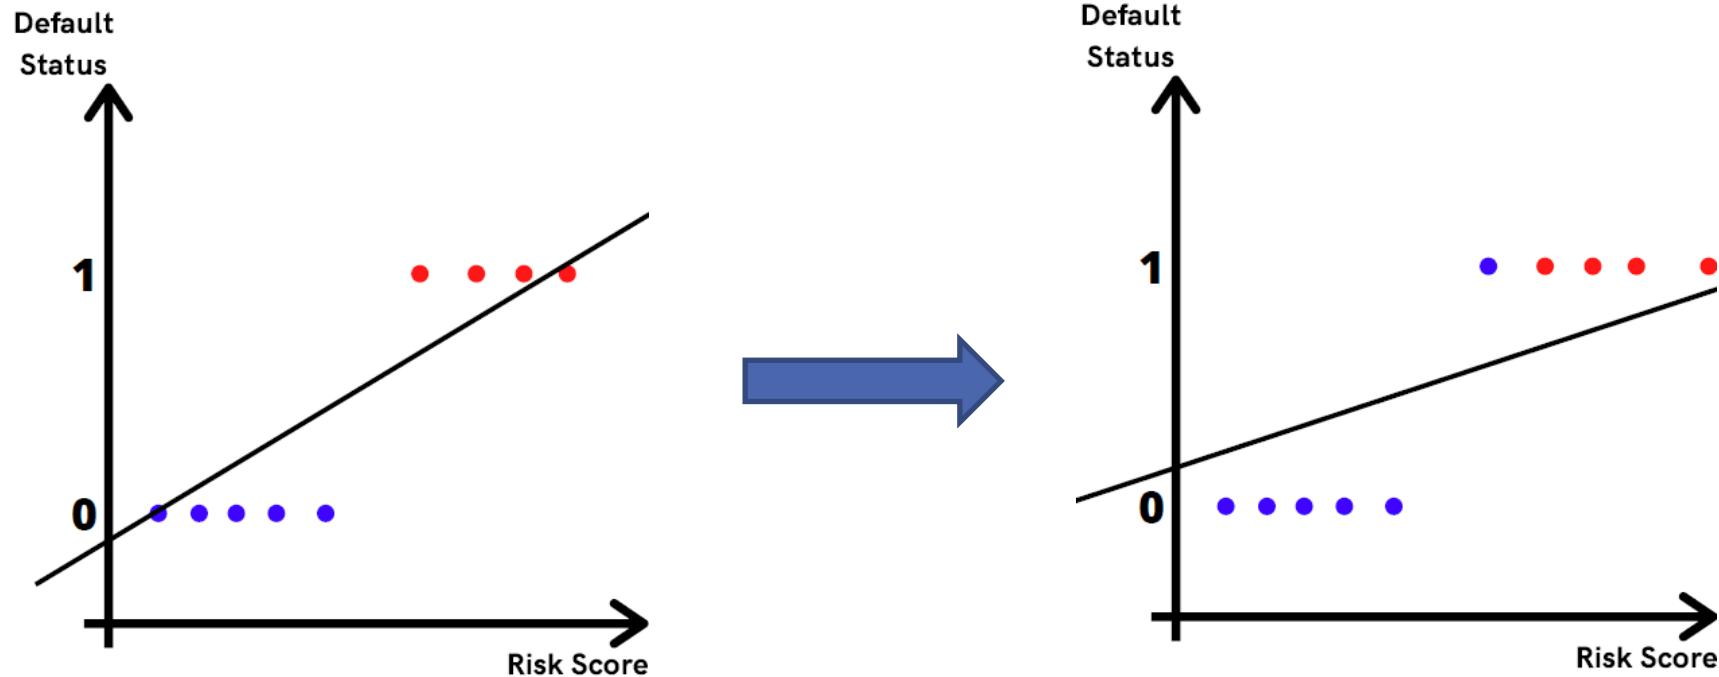


- Sigmoid or Logistic function is an S-shaped curve that is bounded by the interval  $[0,1]$
- The sigmoid function gives the probability that  $Y$  belongs to a particular class

$$p(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

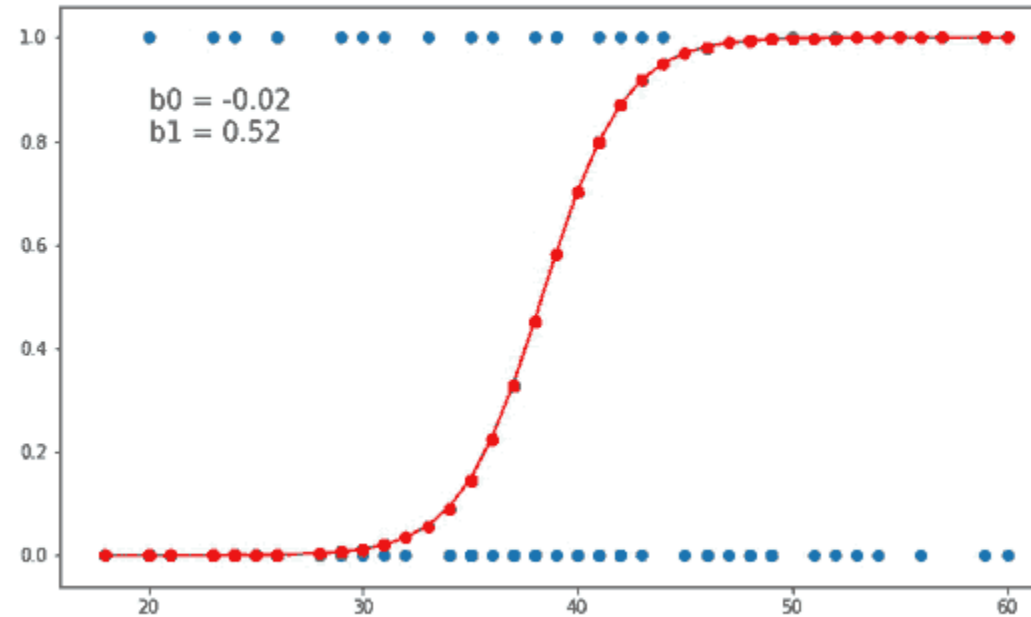
- Whenever  $p(x) \geq 0.5$ ; resulting  $y = 1$
- Whenever  $p(x) < 0.5$ ; resulting  $y = 0$

# Why not using Linear Regression?



- It seems like we can also use Linear Regression to solve classification problems by assigning a threshold value
- But it gets a problem when the linear regression value is greater than 1 or less than 0 will have no meaning value for the classification
- Also, the model tends to make a misclassification when it gets new data, especially outliers

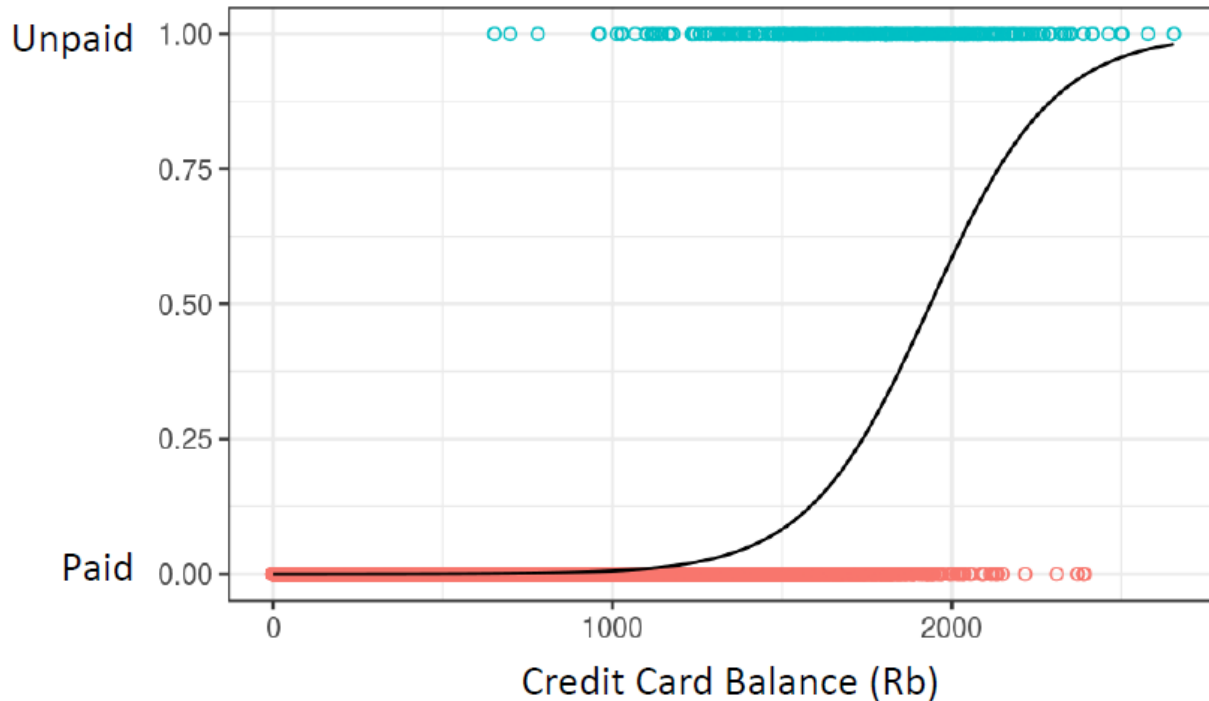
# How Logistic Regression Works using ML?



Machine Learning uses iteration to find the best model for the sigmoid function to suit with the data

# Parameter Vs. Hyperparameter

## Parameter Definition



- A parameter is a variable **estimated from the given data**
- They are required to **make predictions**
- **Not set manually** by the practitioner, meaning it will become the **output**
- Example: coefficients of Logistic Regression

$$p(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$



# Parameter Vs. Hyperparameter

## Hyperparameter Definition

- Hyperparameter is used for **estimating the model parameter**
- They are **specified by the user**
- Tuned for achieving **a specific evaluation metric** such as regularization (decrease probability of overfitting), model adjustment, making the calculation simpler, etc.
- You can learn more about hyperparameter for logistic regression [here](#)

Examples:

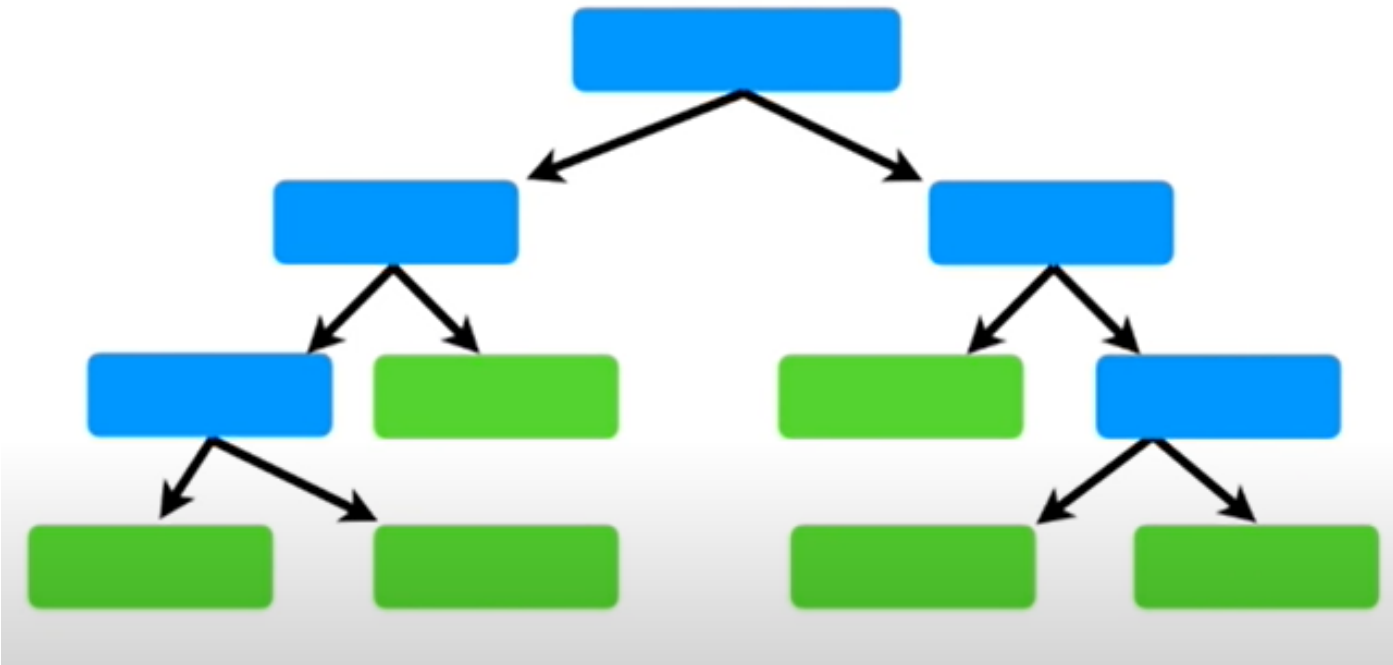
- fit\_intercept** is a Boolean input that decides whether to include the intercept or not

fit_intercept	Function
TRUE	$A = \beta_0 + \beta_1 x$
FALSE	$A = \beta_1 x$

- multi\_class** is a string that decides the approach to use for handling multiple classes ('ovr' by default, 'multinomial', 'auto')
- max\_iter** is an integer (100 by default) that defines the maximum number of iterations by the solver during model fitting

# Supervised Learning: Random Forest

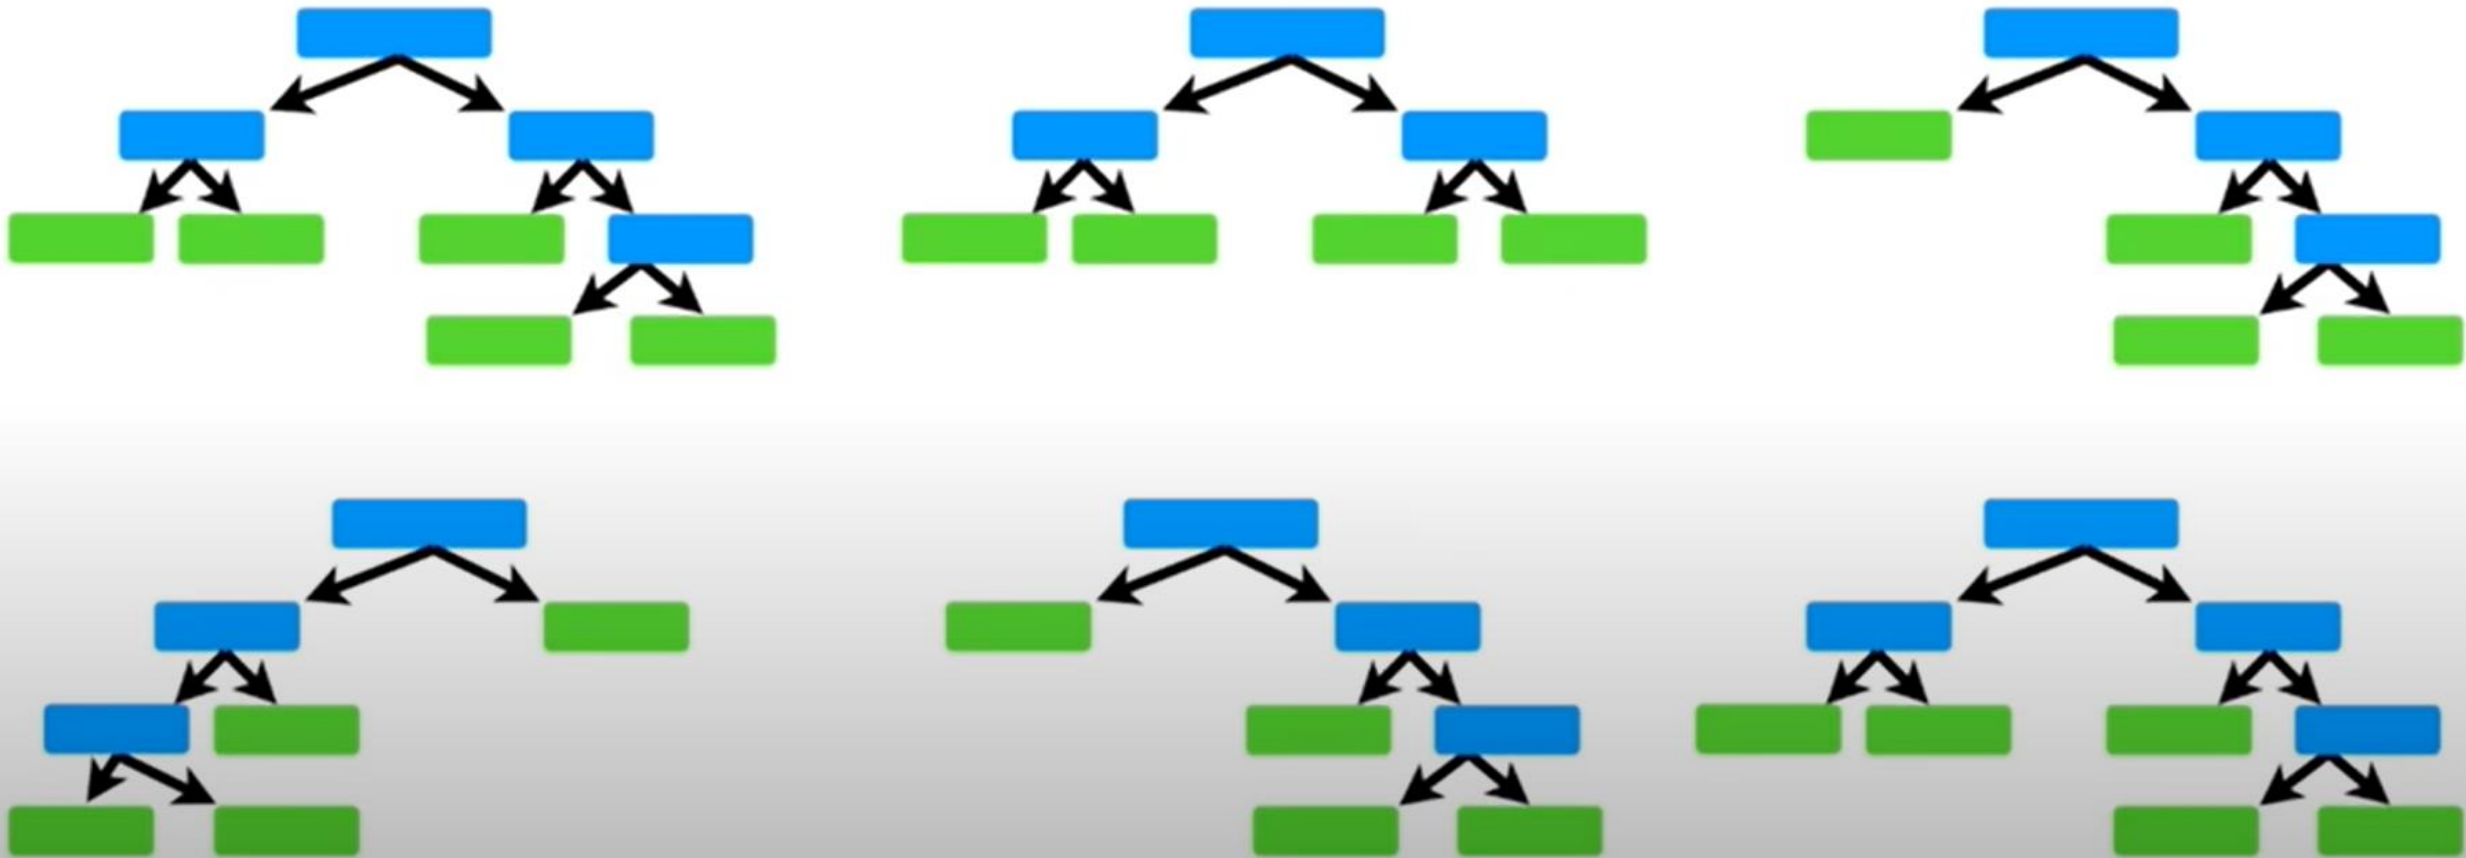
# Why Decision Tree Rarely Used for Classification



- Decision Trees are easy to build, easy to use, and easy to interpret, but it has a bad **inaccuracy**
- This is due to Decision Trees is **inflexible** when it comes to **classifying new samples**

# Why Decision Tree Rarely Used for Classification

Thus, we are using **Random Forest** by combining simple decision trees with flexibility resulting in a vast improvement in accuracy



# How Random Forest Works

Original Dataset

Chest Pain	Good Blood Circ.	Blocked Arteries	Weight	Heart Disease
No	No	No	125	No
Yes	Yes	Yes	180	Yes
Yes	Yes	No	210	No
Yes	No	Yes	167	Yes



Bootstrapped Dataset

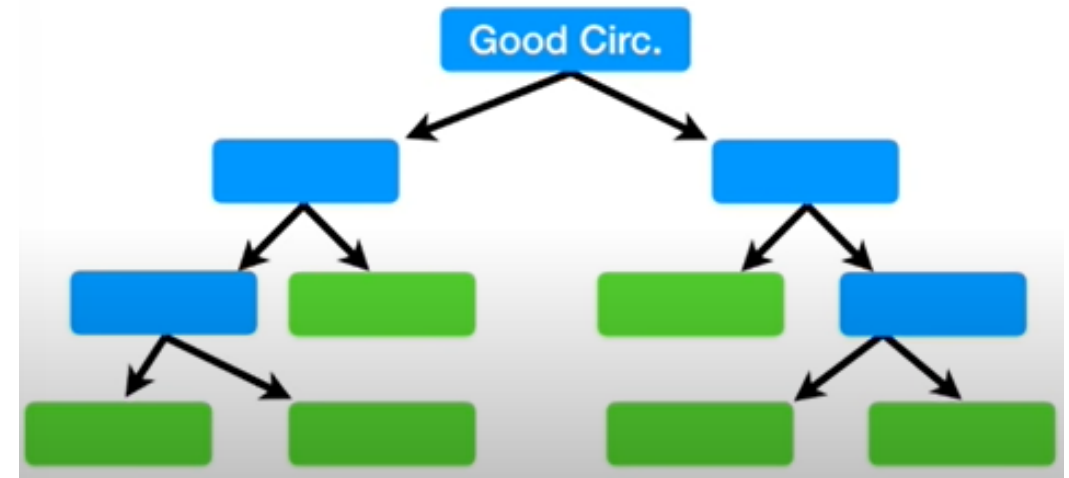
Chest Pain	Good Blood Circ.	Blocked Arteries	Weight	Heart Disease
Yes	Yes	Yes	180	Yes
No	No	No	125	No
Yes	No	Yes	167	Yes
Yes	No	Yes	167	Yes

- Imagine that you have the dataset on the left
- Create a “bootstrapped” dataset by randomly select samples from original (it can be picked more than once)

# How Random Forest Works

Bootstrapped Dataset

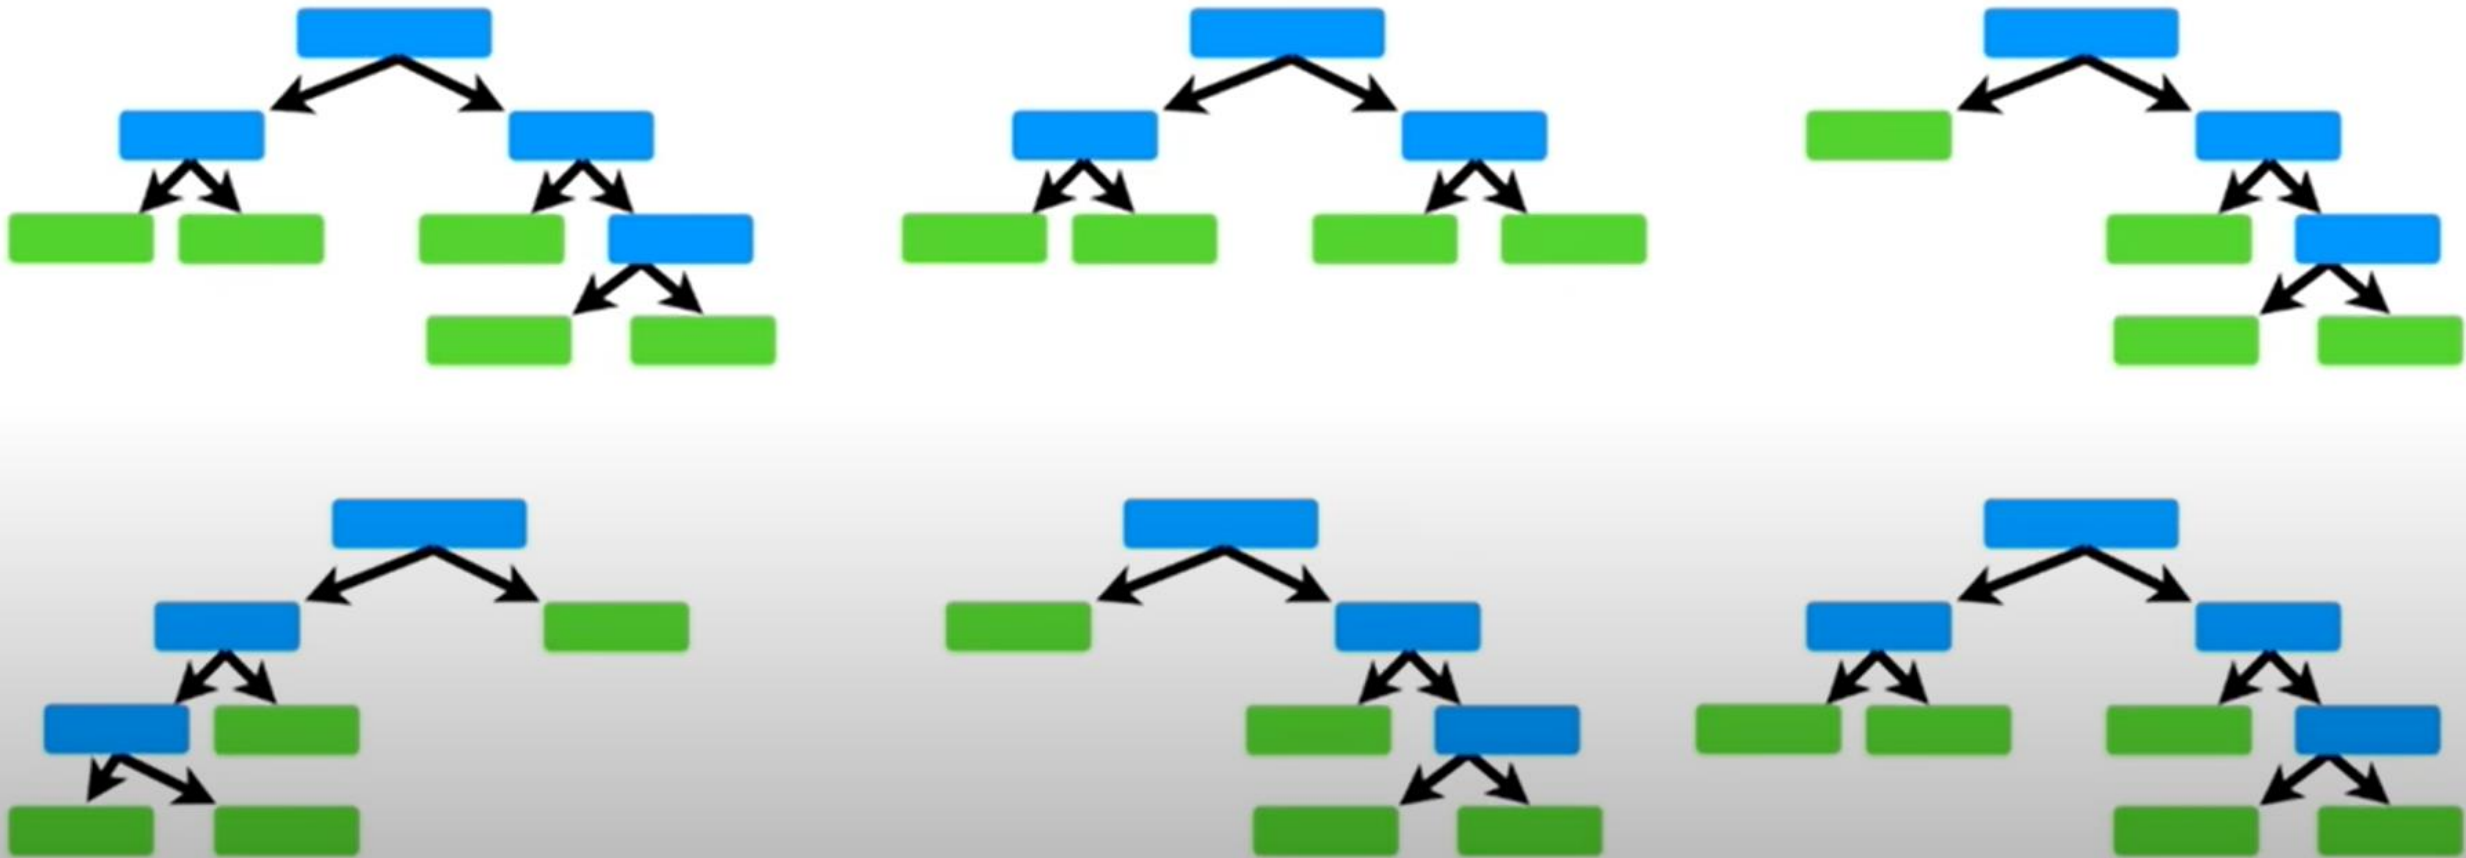
Chest Pain	Good Blood Circ.	Blocked Arteries	Weight	Heart Disease
Yes	Yes	Yes	180	Yes
No	No	No	125	No
Yes	No	Yes	167	Yes
Yes	No	Yes	167	Yes



- Then, create a decision tree using the bootstrapped dataset, but only use a random subset of variables at each step
- The number of step is also limited in order to make a simpler decision tree

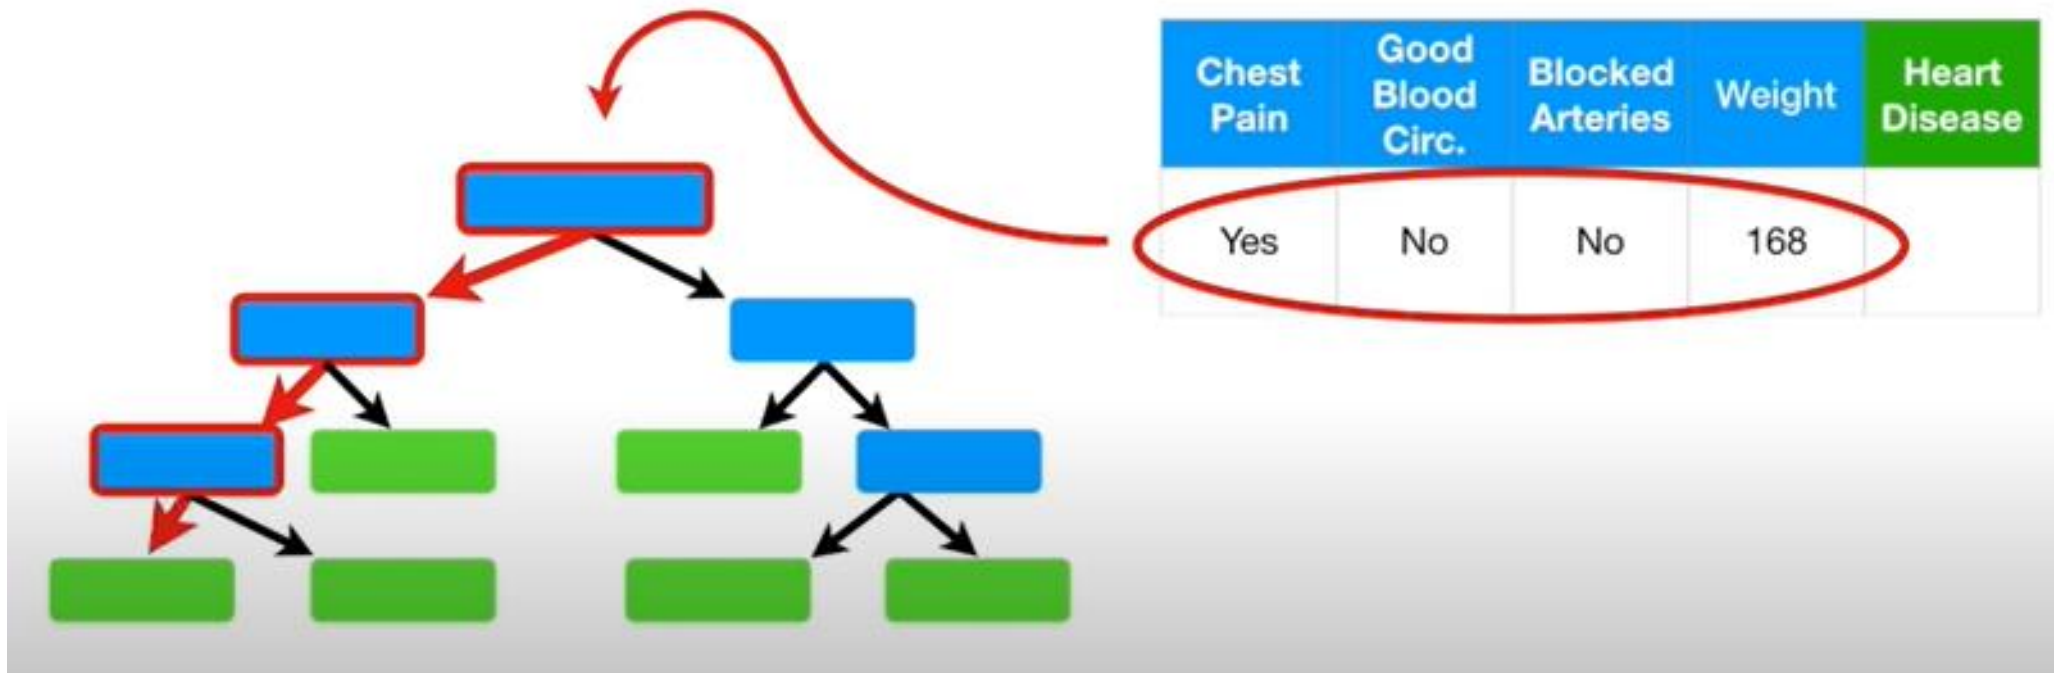
# How Random Forest Works

- Now, repeat the step from the beginning to make multiple decision trees (in default 100 times)



# How Random Forest Works

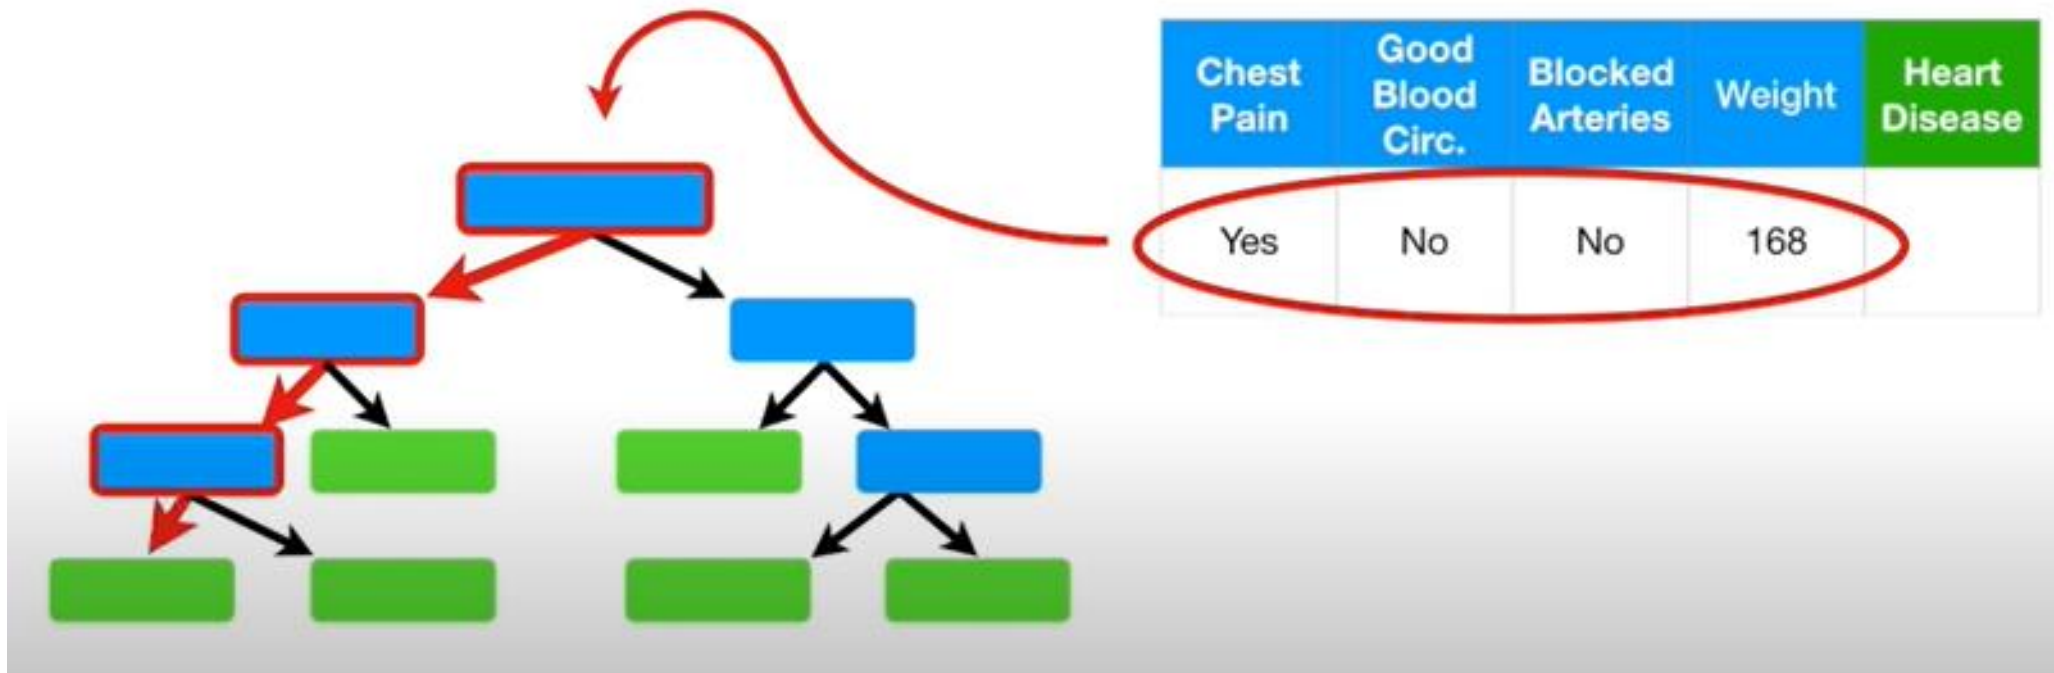
- Now back to the original data to make the predictions by running based on the first decision tree





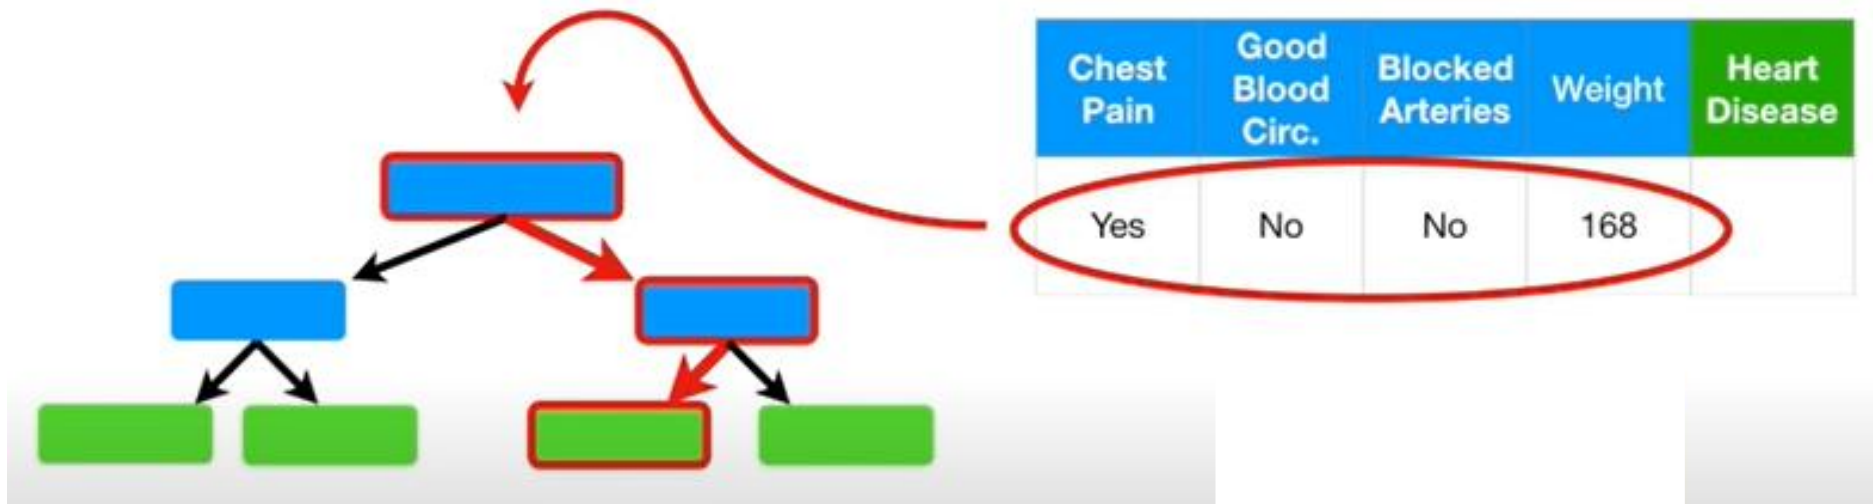
# How Random Forest Works

- Now back to the original data to make the predictions by running based on the first decision tree



# How Random Forest Works

- Then, create another prediction from the second tree and so on



# How Random Forest Works

- It will predict as “YES” when it received the most “YES” votes from all decision trees

Chest Pain	Good Blood Circ.	Blocked Arteries	Weight	Heart Disease
Yes	No	No	168	<b>YES</b>

Heart Disease	
Yes	No
5	1

# Coffee Break

*10:00 - 10:15*





# Technical Details

# Import Package

---

- Import Package Required for Creating Logistic Regression

## Import Package

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
```

# Read Train Data

- Import Data Required for Creating Predictive Model.

## Data For Prediction

```
In [3]: df = pd.read_csv('Data All.csv', low_memory = False)
df.head()
```

Out[3]:

	Customer ID	Branch Code	City	Age	Avg. Annual Income/Month	Balance Q1	NumOfProducts Q1	HasCrCard Q1	ActiveMember Q1	Balance Q2	NumOfProducts Q2	HasCrCard Q2	ActiveMember Q2
0	15565701	1001	Jakarta	29	33000000	0.0	1	1	1	0.0	1	1	0
1	15565878	1005	Jakarta	68	17000000	0.0	2	1	1	0.0	2	1	0
2	15566091	1009	Jakarta	25	12000000	0.0	2	1	0	0.0	2	1	0
3	15566292	1008	Jakarta	42	19000000	0.0	2	1	1	0.0	2	1	0
4	15566312	1009	Jakarta	43	29000000	0.0	2	1	0	0.0	2	1	0

# Data Separation

---

- Separate which data will be considered inside the model
- Separate between text data and numerical data as well

```
In [4]: predictors=df.columns[1:-1]
predictors_onehot = df.columns[1:3]
predictors_num = df.columns[3:-1]
X = df[predictors]
X_onehot = df[predictors_onehot]
X_num = df[predictors_num]
```



# Standard Scaler

- Use Standard Scaler for Numerical Data

```
In [14]: from sklearn.preprocessing import StandardScaler  
pt = StandardScaler()  
X_num = pd.DataFrame(pt.fit_transform(X_num))  
X_num.head()
```

Out[14]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
0	-1.185374	0.663649	-0.95028	-0.869348	0.296502	0.991440	-0.952682	-0.912889	0.296502	-1.046442	-0.954326	-0.967155	0.296502	-1.097147	-0.934539
1	1.580668	-0.680854	-0.95028	0.788861	0.296502	0.991440	-0.952682	0.596812	0.296502	-1.046442	-0.954326	0.424672	0.296502	-1.097147	-0.955811
2	-1.469071	-1.101011	-0.95028	0.788861	0.296502	-1.008634	-0.952682	0.596812	0.296502	-1.046442	-0.954326	0.424672	0.296502	-1.097147	-0.931586
3	-0.263360	-0.512791	-0.95028	0.788861	0.296502	0.991440	-0.952682	0.596812	0.296502	-1.046442	-0.954326	0.424672	0.296502	-1.097147	-0.955811
4	-0.192436	0.327524	-0.95028	0.788861	0.296502	-1.008634	-0.952682	0.596812	0.296502	-1.046442	-0.942169	0.424672	0.296502	0.911455	-0.948084

# One-Hot Encoder

- Use One-Hot Encoder for Categorical Data

```
In [13]: X_onehot = pd.get_dummies(X_onehot, columns = predictors_onehot)
X_onehot.head()
```

Out[13]:

	Branch Code_1001	Branch Code_1002	Branch Code_1003	Branch Code_1004	Branch Code_1005	Branch Code_1006	Branch Code_1007	Branch Code_1008	Branch Code_1009	Branch Code_1011	Branch Code_1012	Branch Code_1013	Branch Code_1014
0	1	0	0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	1	0	0	0	0
3	0	0	0	0	0	0	0	1	0	0	0	0	0
4	0	0	0	0	0	0	0	0	1	0	0	0	0

# Combine Numerical & Categorical Data

- Combine both Numerical & Categorical Data

```
In [16]: X = pd.concat([X_onehot, X_num], axis = 1)  
X.head()
```

Out[16]:

	Branch Code_1001	Branch Code_1002	Branch Code_1003	Branch Code_1004	Branch Code_1005	Branch Code_1006	Branch Code_1007	Branch Code_1008	Branch Code_1009	Branch Code_1011	Branch Code_1012	Branch Code_1013	Branch Code_1014
0	1	0	0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	1	0	0	0	0
3	0	0	0	0	0	0	0	1	0	0	0	0	0
4	0	0	0	0	0	0	0	0	1	0	0	0	0

# Data Test

- Do with the Test Data As Well

## Data for Test

```
In [17]: df_val = pd.read_csv("Data All Test.csv", low_memory=False)
```

```
In [18]: df_val.head()
```

```
Out[18]:
```

	Customer ID	Branch Code	City	Age	Avg. Annual Income/Month	Balance Q2	NumOfProducts Q2	HasCrCard Q2	ActiveMember Q2	Balance Q3	NumOfProducts Q3	HasCrCard Q3	ActiveMember Q3
0	15565701	1001	Jakarta	29	33000000	0.0	1	1	0	0.00	1	1	0
1	15565878	1005	Jakarta	68	17000000	0.0	2	1	0	0.00	2	1	0
2	15566091	1009	Jakarta	25	12000000	0.0	2	1	0	0.00	2	1	0
3	15566292	1008	Jakarta	42	19000000	0.0	2	1	0	0.00	2	1	0
4	15566312	1009	Jakarta	43	29000000	0.0	2	1	0	678905.68	2	1	1

# Import Package for Logistic Regression

- Import package for Logistic Regression

## Logistic Regression

```
In [30]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, stratify=y, random_state=42)
```

```
In [40]: from sklearn.model_selection import RandomizedSearchCV
```

# Building Logistic Regression Model

- Calculate the data in Machine Learning using Logistic Regression

```
In [41]: #Don't Change This
penalty = ['l2']
tol = [0.001, 0.0001, 0.00001]
C = [100.0, 10.0, 1.00, 0.1, 0.01, 0.001]
fit_intercept = [True, False]
intercept_scaling = [1.0, 0.75, 0.5, 0.25]
class_weight = ['balanced', None]
solver = ['newton-cg', 'sag', 'lbfgs', 'saga']
max_iter=[14000]
param_distributions = dict(penalty=penalty, tol=tol, C=C, fit_intercept=fit_intercept, intercept_scaling=intercept_scaling,
                           class_weight=class_weight, solver=solver, max_iter=max_iter)
```

```
In [42]: import time

logreg = LogisticRegression()
grid = RandomizedSearchCV(estimator=logreg, param_distributions = param_distributions , scoring = 'recall', cv = 3, n_jobs=-1)

start_time = time.time()
grid_result = grid.fit(X_train, y_train)
# Summarize results
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
print("Execution time: " + str((time.time() - start_time)) + ' s')

Best: 0.579699 using {'tol': 0.0001, 'solver': 'saga', 'penalty': 'l2', 'max_iter': 14000, 'intercept_scaling': 1.0, 'fit_intercept': True, 'class_weight': 'balanced', 'C': 100.0}
Execution time: 16.35408329963684 s
```

# Calculate the Evaluation Model

## Accuration Test

```
In [43]: y_pred = grid.predict(X_test)
```

```
In [44]: from sklearn import metrics
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
print("Precision:", metrics.precision_score(y_test, y_pred))
print("Recall:", metrics.recall_score(y_test, y_pred))
metrics.completeness_score
```

```
Accuracy: 0.72
Precision: 0.4799588900308325
Recall: 0.5779702970297029
```

```
Out[44]: <function sklearn.metrics.cluster._supervised.completeness_score(labels_true, labels_pred)>
```

```
In [45]: y_pred_val = grid.predict(X_val)
```

```
In [46]: from sklearn import metrics
print("Accuracy:", metrics.accuracy_score(y_val, y_pred_val))
print("Precision:", metrics.precision_score(y_val, y_pred_val))
print("Recall:", metrics.recall_score(y_val, y_pred_val))
metrics.completeness_score
```

```
Accuracy: 0.5435788916809946
Precision: 0.3930835734870317
Recall: 0.5035068290882244
```

- Calculate our results for accuracy, precision, and recall.

Which one we want to consider?  
Why?

# Ishoma

*12:00 - 13:00*







# Case Study: Classification Exercise





Let's Share!





Q & A

*Thank  
you*